

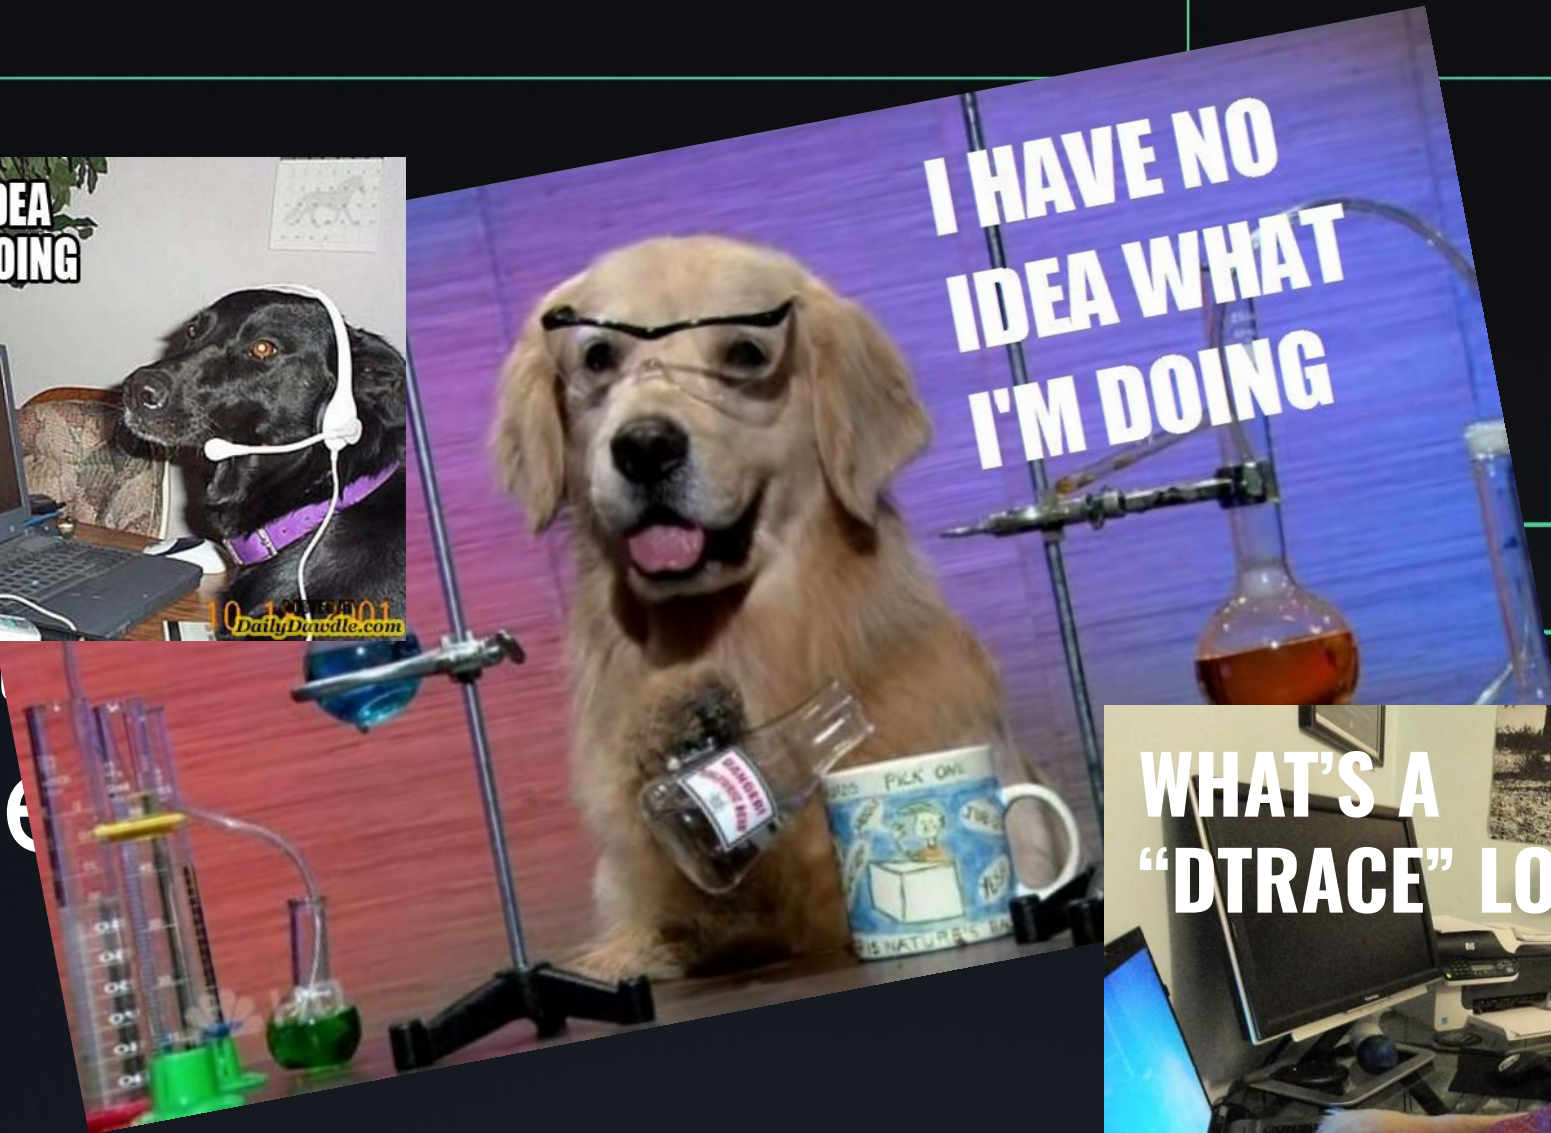



DTracing: Two great tastes...better together?

Eliza Weisman





last
together

Eliza Weisman

WHY I MADE
ANOTHER
LOGGING LIBRARY

A little backstory

- In 2019, I worked at Buoyant, where we were writing a high-performance application-layer proxy in Rust
- We were very **early adopters of the Rust async ecosystem**
- The debugging experience using `log` was **really bad...**

Logs in a **synchronous** system make sense...

DEBUG server: accepted connection from 106.42.126.8:56975

DEBUG server::http: received request

WARN server::http: invalid request headers

TRACE server: closing connection

...but in an **asynchronous system**, they don't

DEBUG server: accepted connection from 106.42.126.8:56975

TRACE server: closing connection

DEBUG server::http: received


DEBUG server: accepted connection from 106.42.126.8:49123

DEBUG server::http: received

DEBUG server: accepted connection from 106.42.126.8:51342

WARN server::http: invalid request headers

TRACE server: closing connection



**THIS IS
NONSENSE!**

Solution: add **context**

DEBUG server{client.addr=106.42.126.8:56975}: accepted connection

TRACE server{client.addr=82.5.70.2:53121}: closing connection

DEBUG server{client.addr=89.56.1.12:55601}:request{path="/posts/tracing"}: received request

DEBUG server{client.addr=111.103.8.9:49123}: accepted connection

DEBUG server{client.addr=106.42.126.8:56975}:request{path="/"}: received request

DEBUG server{client.addr=113.12.37.105:51342}: accepted connection


WARN server{client.addr=106.42.126.8:56975}:request{path="/"}: invalid request headers

TRACE server{client.addr=106.42.126.8:56975}:request{path="/"}: closing connection

Tracing

Structured, application-level diagnostics

- Part of the **Tokio project**
- Data model **inspired by distributed tracing**
- **First-class contexts** (spans)
- A **facade** for modular **subscribers**



Core concepts: **events**

```
tracing::event!(level: Level::INFO, "hello world");
```

```
let question = "life, the universe, and everything";  
tracing::debug!(question, answer = 42);
```

Core concepts: **spans**

```
{  
    let _span = tracing::info_span!("my_span").entered();  
  
    tracing::info!(inside_span = true);  
}
```

```
tracing::info!(inside_span = false);
```

Core concepts: **spans**

```
#[tracing::instrument]
async fn connect(addr: SocketAddr) -> Result<TcpStream> {
    tracing::debug!("connecting...");
    // ...
}
```

So why are we
talking about this
at **dtrace.conf**?

tracing

Unfollow

Application-level tracing for Rust.

[#async](#) [#logging](#) [#metrics](#) [#tracing](#)

Readme

43 Versions

Dependencies

Dependents

Settings

Displaying 1-10 of 8984 reverse dependencies of tracing

tokio

DEPENDS ON ^0.1.29

↓ 250,038,346

An event-driven, non-blocking I/O platform for writing asynchronous I/O backed applications.

hyper

DEPENDS ON ^0.1

↓ 226,683,754

A protective and efficient HTTP library for all.

tokio-util

DEPENDS ON ^0.1.29

↓ 217,146,240

Additional utilities for working with Tokio.

h2

DEPENDS ON ^0.1.35

↓ 204,333,217

An HTTP/2 client and server

async-trait

DEPENDS ON ^0.1.40

↓ 180,550,110

Type erasure for async trait methods

tracing-attributes

DEPENDS ON ^0.1.35

↓ 176,125,310

Procedural macro attributes for automatically instrumenting functions.

Two great tastes...**better together?**

Imagine **one Rust macro** that can:

- Generate a **USDT probe**
- Emit a **text log**
- Located within a **distributed trace**

...all configured at runtime!

tracing and DTrace have some **compatible** **values...**

1. Minimize **disabled** probe effects
2. Don't lie
3. Minimize **enabled** probe effects

Required Methods

✓ `fn enabled(&self, metadata: &Metadata<'_>) -> bool`

[Source](#)

Returns true if a span or event with the specified **metadata** would be recorded.

By default, it is assumed that this filter needs only be evaluated once for each callsite, so it is called by **register_callsite** when each callsite is registered. The result is used to determine if the subscriber is always **interested** or never interested in that callsite. This is intended primarily as an optimization, so that expensive filters (such as those involving string search, et cetera) need not be re-evaluated.

However, if the subscriber's interest in a particular span or event may change, or depends on contexts only determined dynamically at runtime, then the **register_callsite** method should be overridden to return **Interest::sometimes**. In that case, this function will be called every time that span or event occurs.

> `fn new_span(&self, span: &Attributes<'_>) -> Id`

[Source](#)

> `fn record(&self, span: &Id, values: &Record<'_>)`

[Source](#)

> `fn record_follows_from(&self, span: &Id, follows: &Id)`

[Source](#)

> `fn event(&self, event: &Event<'_>)`

[Source](#)

> `fn enter(&self, span: &Id)`

[Source](#)

> `fn exit(&self, span: &Id)`

[Source](#)

```

use ::tracing::__macro_support::Callsite as _;
static __CALLSITE: ::tracing::callsite::DefaultCallsite = {
    static META: ::tracing::Metadata<'static> = {
        ::tracing_core::metadata::Metadata::new(
            "event src/main.rs:2",
            "tracing_demo",
            ::tracing::Level::INFO,
            ::tracing_core::__macro_support::Option::Some("src/main.rs"),
            ::tracing_core::__macro_support::Option::Some(2u32),
            ::tracing_core::__macro_support::Option::Some("tracing_demo"),
            ::tracing_core::field::FieldSet::new(
                &["message"],
                ::tracing_core::callsite::Identifier(&__CALLSITE),
            ),
            ::tracing::metadata::Kind::EVENT,
        )
    };
    ::tracing::callsite::DefaultCallsite::new(&META)
};

let enabled = ::tracing::Level::INFO
    <= ::tracing::level_filters::STATIC_MAX_LEVEL
    && ::tracing::Level::INFO <= ::tracing::level_filters::LevelFilter::current()
    && {
        let interest = __CALLSITE.interest();
        !interest.is_never()
        && ::tracing::__macro_support::__is_enabled(
            __CALLSITE.metadata(),
            interest,
        )
    };

if enabled {
    (|value_set: ::tracing::field::ValueSet| {
        let meta = __CALLSITE.metadata();
        ::tracing::Event::dispatch(meta, &value_set);
    })({
        #[allow(unused_imports)]
        use ::tracing::field::{debug, display, Value};
        let mut iter = __CALLSITE.metadata().fields().iter();
        __CALLSITE
            .metadata()
            .fields()
            .value_set(
                &[
                    (
                        &::tracing::__macro_support::Iterator::next(&mut iter)
                            .expect("FieldSet corrupted (this is a bug)"),
                        ::tracing::__macro_support::Option::Some(
                            &format_args!("hello world") as &dyn Value,
                        ),
                    ),
                ],
            ),
    })
};
} else {
}

```

...but we have some pretty different constraints

- No help from the kernel
- No dynamic binary patching, just **normal Rust code**
- Subscribers and filtering **are in-process**
- Integration with **distributed tracing** (like OpenTelemetry)

Tracing's model for integration **won't work nicely** for USDT...


- A subscriber could turn tracing events into USDT probes...
- ...but this kind of misses the point!
- Both tracing and DTrace rely on **locating instrumentation at the tracepoint**

Solution: **first-party** integration?

- tracing's macros could (optionally) **emit USDT probes** in addition to all the other stuff
- Fortunately, **I'm the first party** and I can just do that (sort of)

Modeling **spans** in DTrace

- A span IDs is just a **u64**
- Generate a **set of probes per span**:
 - **new_span**
 - **enter**
 - **exit**
 - **close**



Thank you

